

DECLARE CURSOR (Transact-SQL)

Applies to: ✓ SQL Server ✓ Azure SQL Database ✓ Azure SQL Managed Instance ✓ SQL database in Microsoft Fabric Preview

Defines the attributes of a Transact-SQL server cursor, such as its scrolling behavior and the query used to build the result set on which the cursor operates. `DECLARE CURSOR` accepts both a syntax based on the ISO standard and a syntax using a set of Transact-SQL extensions.

 [Transact-SQL syntax conventions](#)

☞ Syntax

ISO syntax:

`syntaxsql`

```
DECLARE cursor_name [ INSENSITIVE ] [ SCROLL ] CURSOR
    FOR select_statement
    [ FOR { READ_ONLY | UPDATE [ OF column_name [ , ...n ] ] } ]
[ ; ]
```

Transact-SQL extended syntax:

`syntaxsql`

```
DECLARE cursor_name CURSOR [ LOCAL | GLOBAL ]
    [ FORWARD_ONLY | SCROLL ]
    [ STATIC | KEYSET | DYNAMIC | FAST_FORWARD ]
    [ READ_ONLY | SCROLL_LOCKS | OPTIMISTIC ]
    [ TYPE_WARNING ]
    FOR select_statement
    [ FOR UPDATE [ OF column_name [ , ...n ] ] ]
[ ; ]
```

Arguments

cursor_name

The name of the Transact-SQL server cursor defined. *cursor_name* must conform to the rules for identifiers.

INSENSITIVE

Defines a cursor that makes a temporary copy of the data to be used by the cursor. All requests to the cursor are answered from this temporary table in `tempdb`. Therefore, base table modifications aren't reflected in the data returned by fetches made to this cursor, and this cursor doesn't allow modifications. When ISO syntax is used, if `INSENSITIVE` is omitted, committed deletes and updates made to the underlying tables (by any user) are reflected in subsequent fetches.

SCROLL

Specifies that all fetch options (`FIRST`, `LAST`, `PRIOR`, `NEXT`, `RELATIVE`, `ABSOLUTE`) are available. If `SCROLL` isn't specified in an ISO `DECLARE CURSOR`, `NEXT` is the only fetch option supported. `SCROLL` can't be specified if `FAST_FORWARD` is also specified. If `SCROLL` isn't specified, then only the fetch option `NEXT` is available, and the cursor becomes `FORWARD_ONLY`.

select_statement

A standard `SELECT` statement that defines the result set of the cursor. The keywords `FOR BROWSE`, and `INTO` aren't allowed within *select_statement* of a cursor declaration.

SQL Server implicitly converts the cursor to another type if clauses in *select_statement* conflict with the functionality of the requested cursor type.

READ_ONLY

Prevents updates made through this cursor. The cursor can't be referenced in a `WHERE CURRENT OF` clause in an `UPDATE` or `DELETE` statement. This option overrides the default capability of a cursor to be updated.

UPDATE [OF *column_name* [,...n]]

Defines updatable columns within the cursor. If `OF <column_name> [, <... n>]` is specified, only the columns listed allow modifications. If `UPDATE` is specified without a column list, all columns can be updated.

cursor_name

The name of the Transact-SQL server cursor defined. *cursor_name* must conform to the rules for identifiers.

LOCAL

Specifies that the scope of the cursor is local to the batch, stored procedure, or trigger in which the cursor was created. The cursor name is only valid within this scope. The cursor can be referenced by local cursor variables in the batch, stored procedure, or trigger, or a stored procedure `OUTPUT` parameter. An `OUTPUT` parameter is used to pass the local cursor back to the calling batch, stored procedure, or trigger, which can assign the parameter to a cursor variable to reference the cursor after the stored procedure terminates. The cursor is implicitly deallocated when the batch, stored procedure, or trigger terminates, unless the cursor was passed back in an `OUTPUT` parameter. If it passes back in an `OUTPUT` parameter, the cursor is deallocated when the last variable referencing it is deallocated or goes out of scope.

GLOBAL

Specifies that the scope of the cursor is global to the connection. The cursor name can be referenced in any stored procedure or batch executed by the connection. The cursor is only implicitly deallocated at disconnect.

 **Note**

If neither `GLOBAL` or `LOCAL` is specified, the default is controlled by the setting of the **default to local cursor** database option.

FORWARD_ONLY

Specifies that the cursor can only move forward and be scrolled from the first to the last row. `FETCH NEXT` is the only supported fetch option. All insert, update, and delete statements made by the current user (or committed by other users) that affect rows in the result set, are visible as the rows are fetched. Because the cursor can't be scrolled backward, however, changes made to rows in the database after the row was fetched aren't visible through the cursor. Forward-only cursors are dynamic by default, meaning that all changes are detected as the current row is processed. This provides faster cursor opening and enables the result set to display updates made to the underlying tables. While forward-only cursors don't support backward scrolling, applications can return to the beginning of the result set by closing and reopening the cursor.

If FORWARD_ONLY is specified without the STATIC, KEYSET, or DYNAMIC keywords, the cursor operates as a dynamic cursor. When FORWARD_ONLY or SCROLL aren't specified, FORWARD_ONLY is the default, unless the keywords STATIC, KEYSET, or DYNAMIC are specified. STATIC, KEYSET, and DYNAMIC cursors default to SCROLL. Unlike database APIs such as ODBC and ADO, FORWARD_ONLY is supported with STATIC, KEYSET, and DYNAMIC Transact-SQL cursors.

STATIC

Specifies that the cursor always displays the result set as it was when the cursor was first opened, and makes a temporary copy of the data to be used by the cursor. All requests to the cursor are answered from this temporary table in tempdb. Therefore inserts, updates, and deletes made to base tables aren't reflected in the data returned by fetches made to this cursor, and this cursor doesn't detect changes made to the membership, order, or values of the result set after the cursor is opened. Static cursors might detect their own updates, deletes, and inserts, although they aren't required to do so.

For example, suppose a static cursor fetches a row, and another application then updates that row. If the application refetches the row from the static cursor, the values it sees are unchanged, despite the changes made by the other application. All types of scrolling are supported.

KEYSET

Specifies that the membership and order of rows in the cursor are fixed when the cursor is opened. The set of keys that uniquely identify the rows is built into a table in tempdb known as the *keyset*. This cursor provides functionality between a static and a dynamic cursor in its ability to detect changes. Like a static cursor, it doesn't always detect changes to the membership and order of the result set. Like a dynamic cursor, it does detect changes to the values of rows in the result set.

Keyset-driven cursors are controlled by a set of unique identifiers (keys) known as the keyset. The keys are built from a set of columns that uniquely identify the rows in the result set. The keyset is the set of key values from all the rows returned by the query statement. With keyset-driven cursors, a key is built and saved for each row in the cursor and stored either on the client workstation or on the server. When you access each row, the stored key is used to fetch the current data values from the data source. In a keyset-driven cursor, result set membership is frozen when the keyset is fully populated. Thereafter, additions or updates that affect membership aren't a part of the result set until it reopens.

Changes to data values (made either by the keyset owner or other processes) are visible as the user scrolls through the result set:

- If a row is deleted, an attempt to fetch the row returns an `@@FETCH_STATUS` of -2 because the deleted row appears as a gap in the result set. The key for the row exists in the keyset, but the row no longer exists in the result set.
- Inserts made outside the cursor (by other processes) are visible only if the cursor is closed and reopened. Inserts made from inside the cursor are visible at the end of the result set.
- Updates of key values from outside the cursor resemble a delete of the old row followed by an insert of the new row. The row with the new values isn't visible, and attempts to fetch the row with the old values return an `@@FETCH_STATUS` of -2. The new values are visible if the update is done through the cursor by specifying the `WHERE CURRENT OF` clause.

 **Note**

If the query references at least one table without a unique index, the keyset cursor is converted to a static cursor.

DYNAMIC

Defines a cursor that reflects all data changes made to the rows in its result set as you scroll around the cursor and fetch a new record, regardless of whether the changes occur from inside the cursor or by other users outside the cursor. Therefore all insert, update, and delete statements made by all users are visible through the cursor. The data values, order, and membership of the rows can change on each fetch. The `ABSOLUTE` fetch option isn't supported with dynamic cursors. Updates made outside the cursor aren't visible until they're committed (unless the cursor transaction isolation level is set to `UNCOMMITTED`).

For example, suppose a dynamic cursor fetches two rows, and another application then updates one of those rows and deletes the other. If the dynamic cursor then fetches those rows, it doesn't find the deleted row, but it displays the new values for the updated row.

FAST_FORWARD

Specifies a `FORWARD_ONLY`, `READ_ONLY` cursor with performance optimizations enabled.

`FAST_FORWARD` can't be specified if `SCROLL` or `FOR_UPDATE` is also specified. This type of cursor

doesn't allow data modifications from inside the cursor.

(!) Note

Both `FAST_FORWARD` and `FORWARD_ONLY` can be used in the same `DECLARE CURSOR` statement.

READ_ONLY

Prevents updates made through this cursor. The cursor can't be referenced in a `WHERE CURRENT OF` clause in an `UPDATE` or `DELETE` statement. This option overrides the default capability of a cursor to be updated.

SCROLL_LOCKS

Specifies that positioned updates or deletes made through the cursor are guaranteed to succeed. SQL Server locks the rows as they are read into the cursor to ensure their availability for later modifications. `SCROLL_LOCKS` can't be specified if `FAST_FORWARD` or `STATIC` is also specified.

OPTIMISTIC

Specifies that positioned updates or deletes made through the cursor don't succeed, if the row was updated since it was read into the cursor. SQL Server doesn't lock rows as they're read into the cursor. It instead uses comparisons of `timestamp` column values, or a checksum value if the table has no `timestamp` column, to determine whether the row was modified after it was read into the cursor.

If the row was modified, the attempted positioned update or delete fails. `OPTIMISTIC` can't be specified if `FAST_FORWARD` is also specified.

If `STATIC` is specified along with the `OPTIMISTIC` cursor argument, the combination of the two are implicitly converted to the equivalent of the combination of using `STATIC` and `READ_ONLY` arguments, or the `STATIC` and `FORWARD_ONLY` arguments.

TYPE_WARNING

Specifies that a warning message is sent to the client when the cursor is implicitly converted from the requested type to another.

No warning is sent to the client when the combination of `OPTIMISTIC` and `STATIC` cursor arguments are used, and the cursor is implicitly converted to the equivalent of a `STATIC READ_ONLY` or `STATIC FORWARD_ONLY` cursor. The conversion to `READ_ONLY` turns into a `FAST_FORWARD` and `READ_ONLY` cursor from a clients' perspective.

select_statement

A standard `SELECT` statement that defines the result set of the cursor. The keywords `COMPUTE`, `COMPUTE BY`, `FOR BROWSE`, and `INTO` aren't allowed within *select_statement* of a cursor declaration.

! Note

You can use a query hint within a cursor declaration. However, if you also use the `FOR UPDATE OF` clause, specify `OPTION (<query_hint>)` after `FOR UPDATE OF`.

SQL Server implicitly converts the cursor to another type if clauses in *select_statement* conflict with the functionality of the requested cursor type.

`FOR UPDATE [OF column_name [,...n]]`

Defines updatable columns within the cursor. If `OF <column_name> [, <... n>]` is supplied, only the columns listed allow modifications. If `UPDATE` is specified without a column list, all columns can be updated, unless the `READ_ONLY` concurrency option was specified.

Remarks

`DECLARE CURSOR` defines the attributes of a Transact-SQL server cursor, such as its scrolling behavior and the query used to build the result set on which the cursor operates. The `OPEN` statement populates the result set, and `FETCH` returns a row from the result set. The `CLOSE` statement releases the current result set associated with the cursor. The `DEALLOCATE` statement releases the resources used by the cursor.

The first form of the `DECLARE CURSOR` statement uses the ISO syntax for declaring cursor behaviors. The second form of `DECLARE CURSOR` uses Transact-SQL extensions that allow you to define cursors using the same cursor types used in the database API cursor functions of ODBC or ADO.

You can't mix the two forms. If you specify the SCROLL or INSENSITIVE keywords before the CURSOR keyword, you can't use any keywords between the CURSOR and FOR <select_statement> keywords. If you specify any keywords between the CURSOR and FOR <select_statement> keywords, you can't specify SCROLL or INSENSITIVE before the CURSOR keyword.

If a DECLARE CURSOR using Transact-SQL syntax doesn't specify READ_ONLY, OPTIMISTIC, or SCROLL_LOCKS, the default is as follows:

- If the SELECT statement doesn't support updates (insufficient permissions, accessing remote tables that don't support updates, and so on), the cursor is READ_ONLY.
- STATIC and FAST_FORWARD cursors default to READ_ONLY.
- DYNAMIC and KEYSET cursors default to OPTIMISTIC.

Cursor names can only be referenced by other Transact-SQL statements. They can't be referenced by database API functions. For example, after declaring a cursor, the cursor name can't be referenced from OLE DB, ODBC, or ADO functions or methods. The cursor rows can't be fetched using the fetch functions or methods of the APIs; the rows can only be fetched by Transact-SQL FETCH statements.

After a cursor is declared, these system stored procedures can be used to determine the characteristics of the cursor.

 Expand table

System stored procedures	Description
sp_cursor_list	Returns a list of cursors currently visible on the connection and their attributes.
sp_describe_cursor	Describes the attributes of a cursor, such as whether it's a forward-only or scrolling cursor.
sp_describe_cursor_columns	Describes the attributes of the columns in the cursor result set.
sp_describe_cursor_tables	Describes the base tables accessed by the cursor.

Variables might be used as part of the *select_statement* that declares a cursor. Cursor variable values don't change after a cursor is declared.

Permissions

Permissions of `DECLARE CURSOR` default to any user that has `SELECT` permissions on the views, tables, and columns used in the cursor.

Limitations

You can't use cursors or triggers on a table with a clustered columnstore index. This restriction doesn't apply to nonclustered columnstore indexes. You can use cursors and triggers on a table with a nonclustered columnstore index.

Examples

A. Use basic cursor and syntax

The result set generated at the opening of this cursor includes all rows and all columns in the table. This cursor can be updated, and all updates and deletes are represented in fetches made against this cursor. `FETCH NEXT` is the only fetch available because the `SCROLL` option isn't specified.

SQL

```
DECLARE vend_cursor CURSOR  
    FOR SELECT * FROM Purchasing.Vendor  
OPEN vend_cursor  
FETCH NEXT FROM vend_cursor;
```

B. Use nested cursors to produce report output

The following example shows how cursors can be nested to produce complex reports. The inner cursor is declared for each vendor.

SQL

```
SET NOCOUNT ON;  
  
DECLARE @vendor_id INT, @vendor_name NVARCHAR(50),  
        @message VARCHAR(80), @product NVARCHAR(50);
```

```

PRINT '----- Vendor Products Report -----';

DECLARE vendor_cursor CURSOR FOR
SELECT VendorID, Name
FROM Purchasing.Vendor
WHERE PreferredVendorStatus = 1
ORDER BY VendorID;

OPEN vendor_cursor

FETCH NEXT FROM vendor_cursor
INTO @vendor_id, @vendor_name

WHILE @@FETCH_STATUS = 0
BEGIN
    PRINT ' '
    SELECT @message = '----- Products From Vendor: ' +
        @vendor_name

    PRINT @message

    -- Declare an inner cursor based
    -- on vendor_id from the outer cursor.

    DECLARE product_cursor CURSOR FOR
    SELECT v.Name
    FROM Purchasing.ProductVendor pv, Production.Product v
    WHERE pv.ProductID = v.ProductID AND
    pv.VendorID = @vendor_id -- Variable value from the outer cursor

    OPEN product_cursor
    FETCH NEXT FROM product_cursor INTO @product

    IF @@FETCH_STATUS <> 0
        PRINT '           <<None>>'

    WHILE @@FETCH_STATUS = 0
    BEGIN

        SELECT @message = '           ' + @product
        PRINT @message
        FETCH NEXT FROM product_cursor INTO @product
        END

    CLOSE product_cursor
    DEALLOCATE product_cursor
    -- Get the next vendor.
    FETCH NEXT FROM vendor_cursor
    INTO @vendor_id, @vendor_name
END

```

```
CLOSE vendor_cursor;
DEALLOCATE vendor_cursor;
```

Related content

- [@@FETCH_STATUS \(Transact-SQL\)](#)
- [CLOSE \(Transact-SQL\)](#)
- [Cursors \(Transact-SQL\)](#)
- [DEALLOCATE \(Transact-SQL\)](#)
- [FETCH \(Transact-SQL\)](#)
- [SELECT \(Transact-SQL\)](#)
- [sp_configure \(Transact-SQL\)](#)

Last updated on 11/23/2024